

PA1_Classification

February 4, 2026

1 INITIALIZERS

```
[1]: ### IMPORTS
import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
[2]: ### A HELPER FUNCTION TO DISPLAY SECTION TITLE
def print_section(title):
    print(f'{"="*60}\n{title}\n{"="*60}')

### SEE ALL COLUMNS
pd.set_option('display.max_columns', None)
```

2 DATA LOADING

```
[3]: ### LOAD DATA
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

3 INITIAL DATA INSPECTION

3.1 GET OVERVIEW

```
[4]: #### TRAIN DATASET
print_section('TRAIN')
train.info()
display(train.head())
display(train.describe())

#### TEST DATASET
print_section('TEST')
test.info()
display(test.head())
display(test.describe())
```

=====

TRAIN

=====

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   p_id             614 non-null    int64  
 1   no_times_pregnant 614 non-null    int64  
 2   glucose_concentration 614 non-null    int64  
 3   blood_pressure     614 non-null    int64  
 4   skin_fold_thickness 614 non-null    int64  
 5   serum_insulin      614 non-null    int64  
 6   bmi               614 non-null    float64 
 7   diabetes_pedigree 614 non-null    float64 
 8   age               614 non-null    int64  
 9   diabetes           614 non-null    int64  
dtypes: float64(2), int64(8)
memory usage: 48.1 KB
```

	p_id	no_times_pregnant	glucose_concentration	blood_pressure	\
0	316	2	112	68	
1	25	11	143	94	
2	710	2	93	64	
3	658	1	120	80	
4	542	3	128	72	

	skin_fold_thickness	serum_insulin	bmi	diabetes_pedigree	age	diabetes
0	22	94	34.1	0.315	26	0
1	33	146	36.6	0.254	51	1
2	32	160	38.0	0.674	23	1
3	48	200	38.9	1.162	41	0

```

4           25          190   32.4           0.549    27           1
              p_id  no_times_pregnant  glucose_concentration  blood_pressure \
count    614.000000            614.000000            614.000000            614.000000
mean     385.773616            3.853420           120.542345            68.765472
std      223.603024            3.358126           31.252286            19.914836
min      1.000000            0.000000            0.000000            0.000000
25%     191.250000            1.000000           99.000000            62.000000
50%     387.000000            3.000000           117.000000            72.000000
75%     572.750000            6.000000           139.000000            80.000000
max     768.000000            17.000000          197.000000            114.000000

              skin_fold_thickness  serum_insulin           bmi  diabetes pedigree \
count        614.000000            614.000000            614.000000            614.000000
mean       20.244300            79.355049           31.909935            0.466342
std        15.886083           117.709950           8.007699            0.331090
min       0.000000            0.000000            0.000000            0.078000
25%      0.000000            0.000000           27.300000            0.240250
50%      23.000000            17.000000           32.000000            0.361000
75%      32.000000           126.000000           36.600000            0.613500
max      63.000000            846.000000           59.400000            2.420000

              age  diabetes
count    614.000000  614.000000
mean     33.325733  0.348534
std      11.929569  0.476895
min      21.000000  0.000000
25%     24.000000  0.000000
50%     29.000000  0.000000
75%     41.000000  1.000000
max     81.000000  1.000000
=====

TEST
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 154 entries, 0 to 153
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   p_id             154 non-null    int64  
 1   no_times_pregnant 154 non-null    int64  
 2   glucose_concentration 154 non-null    int64  
 3   blood_pressure     154 non-null    int64  
 4   skin_fold_thickness 154 non-null    int64  
 5   serum_insulin      154 non-null    int64  
 6   bmi               154 non-null    float64 
 7   diabetes_pedigree 154 non-null    float64 
 8   age               154 non-null    int64  

```

```
dtypes: float64(2), int64(7)
```

```
memory usage: 11.0 KB
```

```
    p_id  no_times_pregnant  glucose_concentration  blood_pressure \
0      437                  12                      140                 85
1      411                   6                      102                 90
2      639                   7                      97                 76
3      213                   7                     179                 95
4      181                   6                     87                 80
```

```
    skin_fold_thickness  serum_insulin   bmi  diabetes pedigree  age
0                  33             0  37.4        0.244     41
1                  39             0  35.7        0.674     28
2                  32            91  40.9        0.871     32
3                  31             0  34.2        0.164     60
4                  0              0  23.2        0.084     32
```

```
    p_id  no_times_pregnant  glucose_concentration  blood_pressure \
count  154.000000          154.000000          154.000000          154.000000
mean   379.422078          3.811688           122.298701          70.461039
std    215.338912          3.425719           34.769480          16.935917
min    4.000000            0.000000           0.000000          0.000000
25%   194.250000          1.000000           97.000000          64.000000
50%   383.000000          3.000000           115.000000         72.000000
75%   583.250000          6.000000           144.750000         78.000000
max   738.000000          14.000000          199.000000          122.000000
```

```
    skin_fold_thickness  serum_insulin   bmi  diabetes pedigree \
count  154.000000          154.000000          154.000000          154.000000
mean   21.701299          81.571429          32.322078          0.493942
std    16.213095          105.178271          7.386724          0.332439
min    0.000000            0.000000           0.000000          0.084000
25%   0.000000            0.000000           28.150000          0.248000
50%   24.500000           52.000000          32.900000          0.411500
75%   33.000000           129.500000          36.200000          0.654750
max   99.000000           474.000000          67.100000          2.137000
```

```
    age
count  154.000000
mean   32.902597
std    11.090106
min    21.000000
25%   24.000000
50%   29.500000
75%   40.000000
max   66.000000
```

3.1.1 REMARKS

- 5 non-zero numeric features (`glucose_concentration`, `blood_pressure`, `skin_fold_thickness`, `serum_insulin`, `bmi`) have invalid zeros.
- No column has any missing values.

3.2 IDENTIFY ISSUES FOR DATA CLEANING

A systematic check on potential data issues.

3.2.1 CHECK DATA TYPES

- Are all the columns in both datasets numeric (`int64` or `float64`)?
- Only numeric columns are used in building ML models.

```
[5]: print_section('CHECK DATA TYPES')

def check_data_types(df, name):
    if df.columns.equals(df.select_dtypes(include=['float64', 'int64']).columns):
        print(f' {name}: All data types are numeric (float64 or int64).')
    else:
        print(f' {name}: There are non-numeric data types.')

check_data_types(train, 'train')
check_data_types(test, 'test')
```

```
=====
CHECK DATA TYPES
=====
train: All data types are numeric (float64 or int64).
test: All data types are numeric (float64 or int64).
```

3.2.2 CHECK DUPLICATES

- Are there any duplicated rows?
- Are there any duplicated id?

```
[6]: print_section('CHECK ROW DUPLICATES')

def check_row_dup(df, name):
    row_dup = df.duplicated().sum()
    if row_dup == 0:
        print(f' {name}: {row_dup} row duplicate found.')
    else:
        print(f' {name}: {row_dup} row duplicates found.')

check_row_dup(train, 'train')
check_row_dup(test, 'test')
```

```
=====
```

```
CHECK ROW DUPLICATES
```

```
=====
```

```
train: 0 row duplicate found.  
test: 0 row duplicate found.
```

```
[7]: print_section('CHECK ID DUPLICATES')
```

```
def check_id_dup(df, name):  
    total_entries = len(df)  
    unique_id = df['p_id'].nunique()  
    print(f"Total entries: {total_entries}")  
    print(f"Unique p_id: {unique_id}")  
    if total_entries == unique_id:  
        print(f' {name}: No id duplicate found. 1 unique p_id per patient.\n')  
    else:  
        print(f' {name}: Multiple entries per patient detected!\n')  
  
check_id_dup(train, 'train')  
check_id_dup(test, 'test')
```

```
=====
```

```
CHECK ID DUPLICATES
```

```
=====
```

```
Total entries: 614  
Unique p_id: 614  
train: No id duplicate found. 1 unique p_id per patient.
```

```
Total entries: 154  
Unique p_id: 154  
test: No id duplicate found. 1 unique p_id per patient.
```

3.2.3 CHECK MISSING VALUES

- Are there any columns with missing values?
- If yes, how much do the missing values account for?

```
[8]: print_section('CHECK MISSING VALUES')
```

```
def check_missing(df, name):  
    """ COUNT OF MISSING VALUES  
    miss_cnt = df.isnull().sum()  
  
    """ PERCENTAGE OF MISSING VALUES ROUNDED TO 2 DP  
    miss_pct = (miss_cnt / len(df) * 100).round(2)  
  
    """ DF CONTAINING COLUMNS WITH MISSING VALUES
```

```

miss_df = pd.DataFrame({'missing_count': miss_cnt[miss_cnt>0],
                       'missing_percentage': miss_pct[miss_pct>0]}) \
    .sort_values('missing_count', ascending=False)

if len(miss_df) > 0:
    print('\nColumns with missing values:')
    print(miss_df)
    print(f"\n {name}: Missing values detected.")
    print('Remove columns with % missing more than threshold in DATA')
    ↵CLEANING.')
else:
    print(miss_cnt)
    print(f" {name}: No missing values detected.\n")

check_missing(train, 'train')
check_missing(test, 'test')

```

=====

CHECK MISSING VALUES

=====

```

p_id          0
no_times_pregnant  0
glucose_concentration  0
blood_pressure      0
skin_fold_thickness 0
serum_insulin       0
bmi               0
diabetes_pedigree   0
age               0
diabetes           0
dtype: int64
train: No missing values detected.

```

```

p_id          0
no_times_pregnant  0
glucose_concentration  0
blood_pressure      0
skin_fold_thickness 0
serum_insulin       0
bmi               0
diabetes_pedigree   0
age               0
dtype: int64
test: No missing values detected.

```

3.2.4 CHECK INVALID VALUES

- Are there any columns with invalid values, i.e. values that are unexpected based on domain understanding?
- From the overview earlier, 5 numerical non-zero columns have invalid zeros.

```
[9]: ### 1. DEFINE COLUMNS THAT CAN BE ZERO BUT NOT NEGATIVE
test_zero_cols = ['no_times_pregnant', 'diabetes pedigree']
train_zero_cols = test_zero_cols + ['diabetes']

### 2. DEFINE COLUMNS THAT MUST BE POSITIVE (NOT ZERO/NEGATIVE)
pos_cols = [col for col in train.columns if col not in train_zero_cols]

### 3. UDF TO IDENTIFY INVALID (ZERO/NEGATIVE) VALUES
def check_invalid(df, zero_cols, pos_cols, name):
    print_section(f"CHECK INVALID VALUES - {name}")

    ### COLUMNS THAT CAN BE ZERO BUT NOT NEGATIVE
    neg_only_cnt = (df[zero_cols] < 0).sum()
    neg_only_pct = (neg_only_cnt / len(df) * 100).round(1)
    zero_df = pd.DataFrame({
        "negatives_count": neg_only_cnt,
        "negatives_percent": neg_only_pct
    })
    print("1. COLUMNS THAT CAN BE ZERO BUT NOT NEGATIVE (CHECK NEGATIVE)")
    display(zero_df)

    ### COLUMNS THAT MUST BE POSITIVE (NOT ZERO/NEGATIVE)
    zero_cnt = (df[pos_cols] == 0).sum()
    zero_pct = (zero_cnt / len(df) * 100).round(1)
    neg_cnt = (df[pos_cols] < 0).sum()
    neg_pct = (neg_cnt / len(df) * 100).round(1)
    pos_df = pd.DataFrame({
        "zeros_count": zero_cnt,
        "zeros_percent": zero_pct,
        "negatives_count": neg_cnt,
        "negatives_percent": neg_pct
    })
    print("2. COLUMNS THAT MUST BE POSITIVE (CHECK ZERO/NEGATIVE)")
    display(pos_df)

    return zero_df, pos_df

### 4. OUTPUT
train_zero, train_pos = check_invalid(train, train_zero_cols, pos_cols, 'TRAIN')
test_zero, test_pos = check_invalid(test, test_zero_cols, pos_cols, 'TEST')
```

=====

CHECK INVALID VALUES - TRAIN

```
=====
1. COLUMNS THAT CAN BE ZERO BUT NOT NEGATIVE (CHECK NEGATIVE)
```

	negatives_count	negatives_percent
no_times_pregnant	0	0.0
diabetes_pedigree	0	0.0
diabetes	0	0.0

```
2. COLUMNS THAT MUST BE POSITIVE (CHECK ZERO/NEGATIVE)
```

	zeros_count	zeros_percent	negatives_count	\
p_id	0	0.0	0	
glucose_concentration	4	0.7	0	
blood_pressure	31	5.0	0	
skin_fold_thickness	187	30.5	0	
serum_insulin	304	49.5	0	
bmi	10	1.6	0	
age	0	0.0	0	

	negatives_percent
p_id	0.0
glucose_concentration	0.0
blood_pressure	0.0
skin_fold_thickness	0.0
serum_insulin	0.0
bmi	0.0
age	0.0

```
=====
CHECK INVALID VALUES - TEST
=====
```

```
1. COLUMNS THAT CAN BE ZERO BUT NOT NEGATIVE (CHECK NEGATIVE)
```

	negatives_count	negatives_percent
no_times_pregnant	0	0.0
diabetes_pedigree	0	0.0

```
2. COLUMNS THAT MUST BE POSITIVE (CHECK ZERO/NEGATIVE)
```

	zeros_count	zeros_percent	negatives_count	\
p_id	0	0.0	0	
glucose_concentration	1	0.6	0	
blood_pressure	4	2.6	0	
skin_fold_thickness	40	26.0	0	
serum_insulin	70	45.5	0	
bmi	1	0.6	0	
age	0	0.0	0	

	negatives_percent
p_id	0.0
glucose_concentration	0.0

blood_pressure	0.0
skin_fold_thickness	0.0
serum_insulin	0.0
bmi	0.0
age	0.0

REMARKS

- Both train and test datasets have 5 non-zero columns with invalid zero values.
- We shall treat the invalid values as “missing values”.
- In DATA CLEANING, the columns with “missing values (invalid values)” above a percentage threshold will be dropped, whereas the others below the threshold will be imputed using median.
- Invalid value threshold: 30%
- Columns to impute: glucose_concentration, blood_pressure, bmi
- Columns to drop: skin_fold_thickness, serum_insulin
- `skin_fold_thickness` is decided to be dropped.
 - It has significant invalid percent (26% and 30.5%).
 - It is a proxy for body fat, which is sufficiently represented by `bmi`.
 - Imputing will likely introduce bias for marginal gain in model performance.

3.2.5 CHECK OUTLIERS

- Methods: IQR & box plot.
- TRAIN dataset is used for the bounds for both TRAIN dataset & TEST dataset.

```
[10]: ### DEFINE FEATURES (EXCLUDE ID & TARGET)
feature_cols = [c for c in train.columns if c not in ['p_id', 'diabetes']]
print(f"Feature columns for outlier detection: {feature_cols}")
print(f"\nTotal: {len(feature_cols)}")

### UDF TO COMPUTE IQR BOUNDARIES
# For each column in col, compute IQR-based lower and upper bounds.
# Returns (lower, upper) where lower = Q1 - 1.5*IQR and upper = Q3 + 1.5*IQR.
def compute_iqr_bounds(df, col):
    s = df[col]
    Q1 = s.quantile(0.25)
    Q3 = s.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    return lower, upper

### UDF TO COUNT OUTLIERS BEFORE CAPPING
def outlier_counts_by_bounds(df, cols, bounds_source='train'):
    counts = {}
    ### FOR EACH COLUMN IN COLS
    for col in cols:
```

```

s = df[col]
### COMPUTE IQR BOUNDS
lower, upper = compute_iqr_bounds(df, col)
### COUNT OUTLIERS OUTSIDE THOSE BOUNDS
counts[col] = int(((s < lower) | (s > upper)).sum())
return pd.Series(counts)

```

Feature columns for outlier detection: ['no_times_pregnant',
'glucose_concentration', 'blood_pressure', 'skin_fold_thickness',
'serum_insulin', 'bmi', 'diabetes pedigree', 'age']

Total: 8

```

[11]: ### PREP THE SUBPLOT GRID
n_featureCols = len(feature_cols)
ncols = 2 # TRAIN / TEST
nrows = n_featureCols # ONE ROW PER FEATURE
fig, axes = plt.subplots(nrows, ncols, figsize=(10, nrows*3))

print_section('CHECK OUTLIERS (OUTSIDE IQR BOUNDS OF TRAIN)')

### GET OUTLIER COUNTS
before_train_counts = outlier_counts_by_bounds(train, feature_cols)
before_test_counts = outlier_counts_by_bounds(test, feature_cols)

### FOR EACH COLUMN IN TRAIN/TEST, PLOT THE BOXPLOT & SHOW OUTLIER COUNT
for i, c in enumerate(feature_cols):
    ### TRAIN BOXPLOT (LEFT COLUMN)
    axes[i, 0].boxplot(train[c], vert=False)
    axes[i, 0].set_title(f'\nTrain - {c}')
    ### RED IF OUTLIERS; GREEN IF NO OUTLIERS
    train_color = 'green' if before_train_counts[c] == 0 else 'red'
    axes[i, 0].text(1.05, 1, f"Outliers: {before_train_counts[c]}",
                    transform=axes[i, 0].transAxes, fontsize=10,
                    color=train_color, va='center')

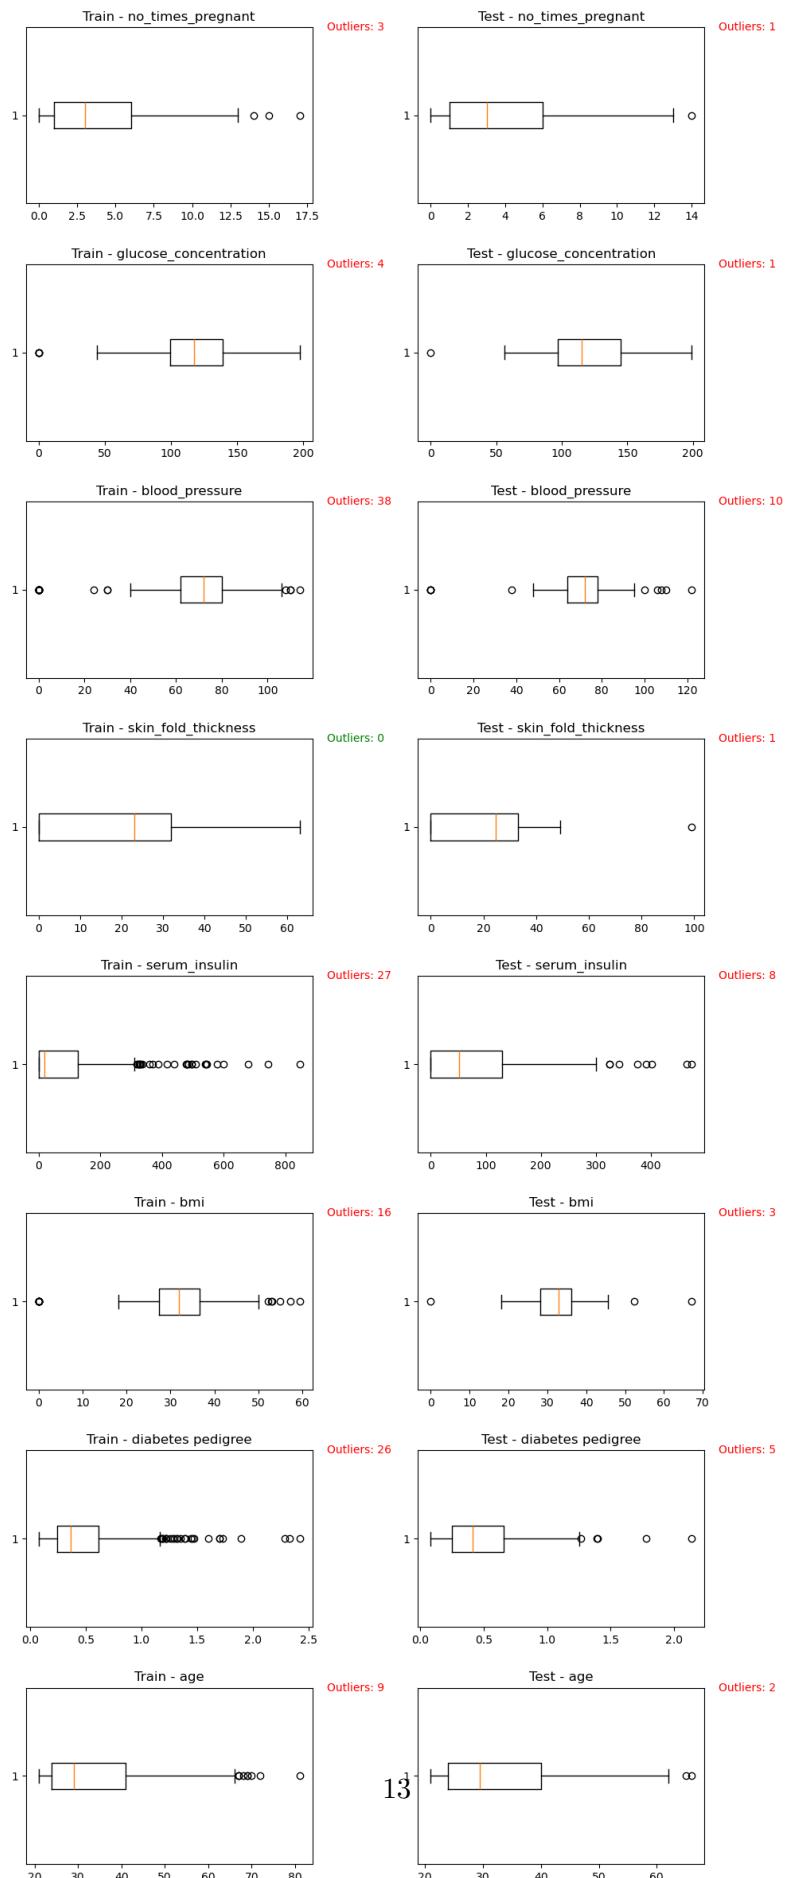
    axes[i, 1].boxplot(test[c], vert=False)
    axes[i, 1].set_title(f'\nTest - {c}')
    ### RED IF OUTLIERS; GREEN IF NO OUTLIERS
    test_color = 'green' if before_test_counts[c] == 0 else 'red'
    axes[i, 1].text(1.05, 1, f"Outliers: {before_test_counts[c]}",
                    transform=axes[i, 1].transAxes, fontsize=10,
                    color=test_color, va='center')

plt.tight_layout()
plt.show()

```

CHECK OUTLIERS (OUTSIDE IQR BOUNDS OF TRAIN)

=====



REMARKS

- Both datasets have outliers, except `skin_fold_thickness` in train.
- We shall ignore the outlier plots in `skin_fold_thickness` and `insulin_serum` which will be dropped later in DATA CLEANING.
- For other outliers, they will be **intentionally left unmodified** during DATA CLEANING as they represent legitimate, extreme patient medical characteristics, rather than systematic data errors, due to the small percentage they account for.

4 DATA CLEANING

- From INITIAL DATA INSPECTION, there are no invalid data types, duplicates, or missing values.
- Columns with invalid values will be imputed or dropped.
- Outliers will be left unmodified.

4.1 HANDLE INVALID VALUES

From INITIAL DATA INSPECTION, - Invalid value threshold: 30% - Columns to impute: `glucose_concentration`, `blood_pressure`, `bmi` - Columns to drop: `skin_fold_thickness`, `serum_insulin` - `skin_fold_thickness` is decided to be dropped. - It has significant invalid percent (26% and 30.5%). - It is a proxy for body fat, which is sufficiently represented by `bmi`. - Imputing will likely introduce bias for marginal gain in model performance.

```
[12]: print_section('HANDLE INVALID VALUES')

### DEFINE COLUMNS TO DROP/IMPUTE
cols_drop = ['skin_fold_thickness', 'serum_insulin']
cols_imp = ['glucose_concentration', 'blood_pressure', 'bmi']

### DROP COLUMNS & CREATE COPIES FOR IMPUTATION
print(f"Dropping columns {cols_drop}...\n")
train_imp = train.drop(cols_drop, axis=1).copy()
test_imp = test.drop(cols_drop, axis=1).copy()

### SANITY CHECK
print(f"Train: {train.shape} -> {train_imp.shape}")
print(f"Test: {test.shape} -> {test_imp.shape}")

### IMPUTE ZEROS WITH TRAIN'S MEDIAN VALUES
print(f"\nImputing zeros in {cols_imp}...\n")
for col in cols_imp:
    ### IMPUTE
    median_val = train_imp[col].median()
    train_imp[col] = train_imp[col].replace(0, median_val)
```

```

test_imp[col] = test_imp[col].replace(0, median_val)

### SANITY CHECK
zeros_train = (train_imp[col] == 0).sum()
zeros_test = (test_imp[col] == 0).sum()
print(f"{col}: train zeros = {zeros_train}, test zeros = {zeros_test}")

### LAST CHECK
print("\nColumn Stats")
display(train_imp.describe())
display(test_imp.describe())

```

=====

HANDLE INVALID VALUES

=====

Dropping columns ['skin_fold_thickness', 'serum_insulin']...

Train: (614, 10) -> (614, 8)
Test: (154, 9) -> (154, 7)

Imputing zeros in ['glucose_concentration', 'blood_pressure', 'bmi']...

glucose_concentration: train zeros = 0, test zeros = 0
blood_pressure: train zeros = 0, test zeros = 0
bmi: train zeros = 0, test zeros = 0

Column Stats

	p_id	no_times_pregnant	glucose_concentration	blood_pressure	\
count	614.000000	614.000000	614.000000	614.000000	
mean	385.773616	3.853420	121.304560	72.400651	
std	223.603024	3.358126	29.688213	12.031595	
min	1.000000	0.000000	44.000000	24.000000	
25%	191.250000	1.000000	100.000000	64.000000	
50%	387.000000	3.000000	117.000000	72.000000	
75%	572.750000	6.000000	139.000000	80.000000	
max	768.000000	17.000000	197.000000	114.000000	

	bmi	diabetes	pedigree	age	diabetes
count	614.000000	614.000000	614.000000	614.000000	
mean	32.431107	0.466342	33.325733	0.348534	
std	6.873171	0.331090	11.929569	0.476895	
min	18.200000	0.078000	21.000000	0.000000	
25%	27.500000	0.240250	24.000000	0.000000	
50%	32.000000	0.361000	29.000000	0.000000	
75%	36.600000	0.613500	41.000000	1.000000	
max	59.400000	2.420000	81.000000	1.000000	

p_id	no_times_pregnant	glucose_concentration	blood_pressure	\
------	-------------------	-----------------------	----------------	---

count	154.000000	154.000000	154.000000	154.000000
mean	379.422078	3.811688	123.058442	72.331169
std	215.338912	3.425719	33.328085	12.392310
min	4.000000	0.000000	56.000000	38.000000
25%	194.250000	1.000000	97.500000	66.000000
50%	383.000000	3.000000	115.500000	72.000000
75%	583.250000	6.000000	144.750000	78.000000
max	738.000000	14.000000	199.000000	122.000000

	bmi	diabetes	pedigree	age
count	154.000000	154.000000	154.000000	154.000000
mean	32.52987	0.493942	32.902597	
std	6.90599	0.332439	11.090106	
min	18.20000	0.084000	21.000000	
25%	28.37500	0.248000	24.000000	
50%	32.90000	0.411500	29.500000	
75%	36.20000	0.654750	40.000000	
max	67.10000	2.137000	66.000000	

4.2 LATEST CLEANED DATA

```
[13]: train_clean = train_imp.copy()
test_clean = test_imp.copy()
feature_cols_clean = [c for c in train_clean.columns if c not in ['p_id', ↴
    'diabetes']]
print(f"Feature Columns After Cleaning: {feature_cols_clean}\n")

print_section('TRAIN_CLEAN')
train_clean.info()
display(train_clean.head())
print_section('TEST_CLEAN')
test_clean.info()
display(test_clean.head())

print(''
    DATA CLEANING COMPLETED!
    - TRAIN DATASET IS READY FOR EDA.
    - TEST DATASET IS READY FOR FEATURE ENGINEERING.
    '')
```

Feature Columns After Cleaning: ['no_times_pregnant', 'glucose_concentration',
'blood_pressure', 'bmi', 'diabetes pedigree', 'age']

```
=====
TRAIN_CLEAN
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
```

```

Data columns (total 8 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   p_id              614 non-null    int64  
 1   no_times_pregnant 614 non-null    int64  
 2   glucose_concentration 614 non-null  int64  
 3   blood_pressure     614 non-null    int64  
 4   bmi                614 non-null    float64 
 5   diabetes_pedigree 614 non-null    float64 
 6   age                614 non-null    int64  
 7   diabetes           614 non-null    int64  
dtypes: float64(2), int64(6)
memory usage: 38.5 KB

      p_id  no_times_pregnant  glucose_concentration  blood_pressure  bmi  \
0      316                  2                      112             68  34.1
1       25                  11                     143             94  36.6
2      710                  2                      93              64  38.0
3      658                  1                     120             80  38.9
4      542                  3                     128             72  32.4

   diabetes_pedigree  age  diabetes
0          0.315    26      0
1          0.254    51      1
2          0.674    23      1
3          1.162    41      0
4          0.549    27      1
=====

TEST_CLEAN
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 154 entries, 0 to 153
Data columns (total 7 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   p_id              154 non-null    int64  
 1   no_times_pregnant 154 non-null    int64  
 2   glucose_concentration 154 non-null  int64  
 3   blood_pressure     154 non-null    int64  
 4   bmi                154 non-null    float64 
 5   diabetes_pedigree 154 non-null    float64 
 6   age                154 non-null    int64  
dtypes: float64(2), int64(5)
memory usage: 8.6 KB

      p_id  no_times_pregnant  glucose_concentration  blood_pressure  bmi  \
0      437                  12                      140             85  37.4
1      411                   6                      102             90  35.7

```

```

2    639          7          97          76  40.9
3    213          7         179          95  34.2
4    181          6          87          80  23.2

      diabetes pedigree  age
0            0.244    41
1            0.674    28
2            0.871    32
3            0.164    60
4            0.084    32

```

DATA CLEANING COMPLETED!

- TRAIN DATASET IS READY FOR EDA.
- TEST DATASET IS READY FOR FEATURE ENGINEERING.

5 EDA

- EDA is done only on the TRAIN dataset without touching the TEST dataset.
- This is to prevent leaking information and artificially improving ML model performance.

```
[14]: ### SET PLOTTING STYLE
# white background, visible grey gridlines
# all figure sizes to 12 inches wide x 6 inches tall unless specified
sns.set_theme(style='whitegrid', rc={'figure.figsize': (14, 6)})
```

5.1 TARGET VARIABLE ANALYSIS

```
[15]: print_section('TARGET VARIABLE ANALYSIS')

### SHOW COUNT AND PERCENTAGE
diabetes_count = train_clean['diabetes'].value_counts()
diabetes_pct = train_clean['diabetes'].value_counts(normalize=True) * 100

# :.1f and :.2f to round to 1 and 2 decimal place(s) respectively
print('Diabetes Distribution:')
print(f"Non-Diabetic (0): {diabetes_count[0]} ({diabetes_pct[0]:.1f}%)")
print(f"Diabetic (1): {diabetes_count[1]} ({diabetes_pct[1]:.1f}%)")
print(f"\nNon-Diabetic-to-Diabetic Ratio: {diabetes_count[0]/diabetes_count[1]:.2f}:1")

### VISUALIZE TARGET DISTRIBUTION
# Create a figure of two subplots (1 row x 2 cols)
# -----
# /   Axes[0]   /   Axes[1]   /
# -----
fig, axes = plt.subplots(1, 2)
```

```

colors = ['#2ecc71', '#e74c3c']

# Bar Chart for Axes[0]
diabetes_count.plot(kind='bar', ax=axes[0], color=colors, alpha=0.8)
axes[0].set_title('Diabetes Distribution (Count)', fontsize=12, u
    ↪fontweight='bold')
axes[0].set_xlabel('Diabetes (0=No, 1=Yes)', fontsize=10)
axes[0].set_ylabel('Count', fontsize=10)
axes[0].set_xticklabels(['No Diabetes', 'Diabetes'], rotation=0)

# Add count labels on bars (5 counts above current count y)
# va='bottom': the bottom of text box sits above y-coordinate = y + 5
for x, y in enumerate(diabetes_count):
    axes[0].text(x, y + 5, str(y), ha='center', va='bottom')

# Pie Chart for Axes[1]
axes[1].pie(diabetes_count, labels=['No Diabetes', 'Diabetes'],
    colors=colors, autopct='%1.0f%%', startangle=90,
    explode=(0.05, 0.05), shadow=True)
axes[1].set_title('Diabetes Distribution (Percentage)', fontsize=12, u
    ↪fontweight='bold')

plt.tight_layout()
plt.show()

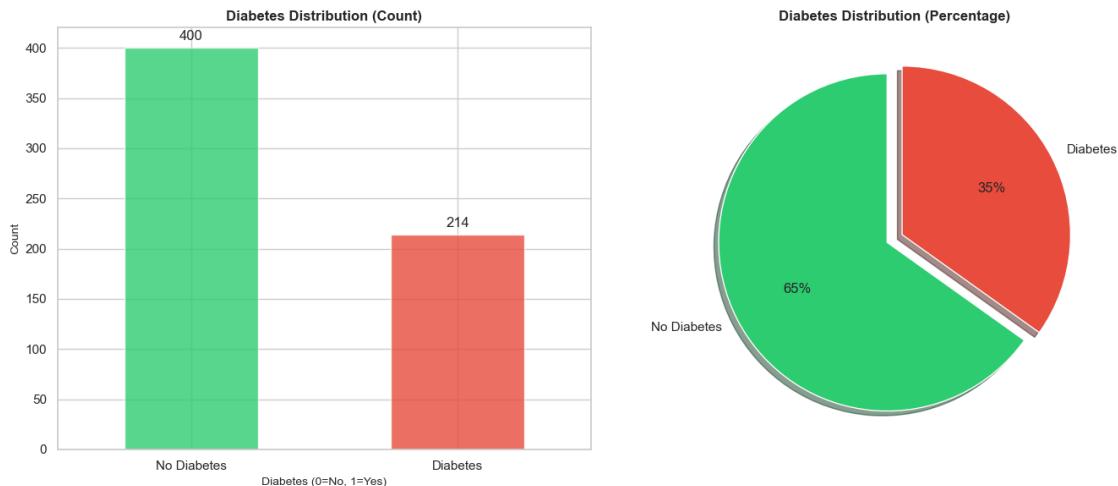
```

TARGET VARIABLE ANALYSIS

Diabetes Distribution:

Non-Diabetic (0): 400 (65.1%)
 Diabetic (1): 214 (34.9%)

Non-Diabetic-to-Diabetic Ratio: 1.87:1



5.1.1 REMARKS

- The train dataset has class imbalance - 1.87 times more non-diabetic than diabetic patients.
- This can make the ML models more biased towards the majority class (non-diabetic) and predict it more frequently.
- Hence, the “accuracy” metric becomes unreliable, and cannot be the sole evaluation metric.
- During model evaluation, **use other metrics (precision, recall, F1-score) alongside accuracy** to give a better picture of a model’s performance.
- <https://www.geeksforgeeks.org/machine-learning/handling-imbalanced-data-for-classification/>

5.2 STATISTICAL SUMMARY BY DIABETES STATUS

```
[16]: print_section('STATISTICAL SUMMARY BY DIABETES')

### SUMMARY STATISTICS GROUPED BY DIABETES STATUS
print('\n--- Non-Diabetic (diabetes=0) ---')
non_diabetic_stat = train_clean[train_clean['diabetes']==0][feature_cols_clean].describe()
display(non_diabetic_stat.round(2))

print('\n--- Diabetic (diabetes=1) ---')
diabetic_stat = train_clean[train_clean['diabetes']==1][feature_cols_clean].describe()
display(diabetic_stat.round(2))

### COMPARE MEDIANS BETWEEN GROUPS
print('\n--- Median Comparison: Diabetic vs Non-Diabetic ---')
median_compare = pd.DataFrame({
    'Non-Diabetic': train_clean[train_clean['diabetes']==0][feature_cols_clean].median(),
    'Diabetic': train_clean[train_clean['diabetes']==1][feature_cols_clean].median(),
})
median_compare['Difference'] = median_compare['Diabetic'] - median_compare['Non-Diabetic']
median_compare['Pct_Change'] = (median_compare['Difference'] / median_compare['Non-Diabetic']) * 100.round(0)
display(median_compare.round(2))
```

=====

STATISTICAL SUMMARY BY DIABETES

=====

--- Non-Diabetic (diabetes=0) ---

	no_times_pregnant	glucose_concentration	blood_pressure	bmi	\
count	400.00	400.00	400.00	400.00	400.00
mean	3.34	110.41	70.78	30.94	
std	3.01	23.78	11.61	6.58	
min	0.00	44.00	24.00	18.20	
25%	1.00	94.00	62.00	25.90	
50%	2.00	108.00	72.00	30.10	
75%	5.00	124.25	78.00	35.30	
max	13.00	197.00	110.00	57.30	

	diabetes	pedigree	age
count	400.00	400.00	
mean	0.42	31.39	
std	0.30	12.00	
min	0.08	21.00	
25%	0.23	23.00	
50%	0.33	27.00	
75%	0.55	37.00	
max	2.33	81.00	

--- Diabetic (diabetes=1) ---

	no_times_pregnant	glucose_concentration	blood_pressure	bmi	\
count	214.00	214.00	214.00	214.00	214.00
mean	4.81	141.67	75.43	35.22	
std	3.74	28.93	12.25	6.54	
min	0.00	78.00	30.00	22.90	
25%	1.25	119.00	68.00	30.50	
50%	4.00	140.00	74.00	34.05	
75%	8.00	164.75	84.00	38.92	
max	17.00	197.00	114.00	59.40	

	diabetes	pedigree	age
count	214.00	214.00	
mean	0.54	36.94	
std	0.37	10.95	
min	0.12	21.00	
25%	0.26	28.00	
50%	0.43	36.00	
75%	0.73	44.75	
max	2.42	70.00	

--- Median Comparison: Diabetic vs Non-Diabetic ---

	Non-Diabetic	Diabetic	Difference	Pct_Change
no_times_pregnant	2.00	4.00	2.00	100.0
glucose_concentration	108.00	140.00	32.00	30.0
blood_pressure	72.00	74.00	2.00	3.0

bmi	30.10	34.05	3.95	13.0
diabetes pedigree	0.33	0.43	0.11	32.0
age	27.00	36.00	9.00	33.0

5.2.1 REMARKS

- Median values are consistently higher in diabetic patients across all six numerical features, indicating a systematic shift in central tendency rather than effects driven by outliers.
- The most pronounced differences are observed in glucose concentration (+30%), age (+33%), number of pregnancies (median increase from 2 to 4), and diabetes pedigree score (+32%). BMI shows a moderate increase (+13%), while blood pressure exhibits only a small median difference (+3%).
- These results are descriptive but do not imply statistical significance.

5.3 FEATURE DISTRIBUTIONS BY DIABETES

```
[17]: print_section('FEATURE DISTRIBUTIONS BY DIABETES')

### SET THE SUBPLOT GRID SHAPE
n_cols = 2
# Count no. of rows of subplots based on feature cols
n_rows = int(np.ceil(len(feature_cols_clean) / n_cols))
# Create the grid 18-inch wide with 4 inches per row
fig, axes = plt.subplots(n_rows, n_cols, figsize=(18, n_rows*4))
# Flatten the 2D array of axes (n_rows, n_cols) to 1D for easy indexing of each subplot as axes(idx)
axes = axes.ravel()

### FOR EVERY COLUMN AND ITS ENUMERATED INDEX IN feature_cols_clean
for idx, col in enumerate(feature_cols_clean):
    ### PLOT FOR NON-DIABETIC
    train_clean[train_clean['diabetes']==0][col].hist(
        bins=20, alpha=0.6, label='No Diabetes',
        color=colors[0], ax=axes[idx], edgecolor='black'
    )
    ### PLOT FOR DIABETIC
    train_clean[train_clean['diabetes']==1][col].hist(
        bins=20, alpha=0.6, label='Diabetes',
        color=colors[1], ax=axes[idx], edgecolor='black'
    )
    ### SET PLOT ANNOTATIONS
    axes[idx].set_title(f'Distribution of {col}', fontsize=12, fontweight='bold')
    axes[idx].set_xlabel(col, fontsize=10)
    axes[idx].set_ylabel('Frequency', fontsize=10)
    axes[idx].legend()
    axes[idx].grid(axis='y', alpha=0.2)
```

```

#### HIDE UNUSED SUBPLOT (IF ANY)
for idx in range(len(feature_cols_clean), len(axes)):
    axes[idx].axis('off')

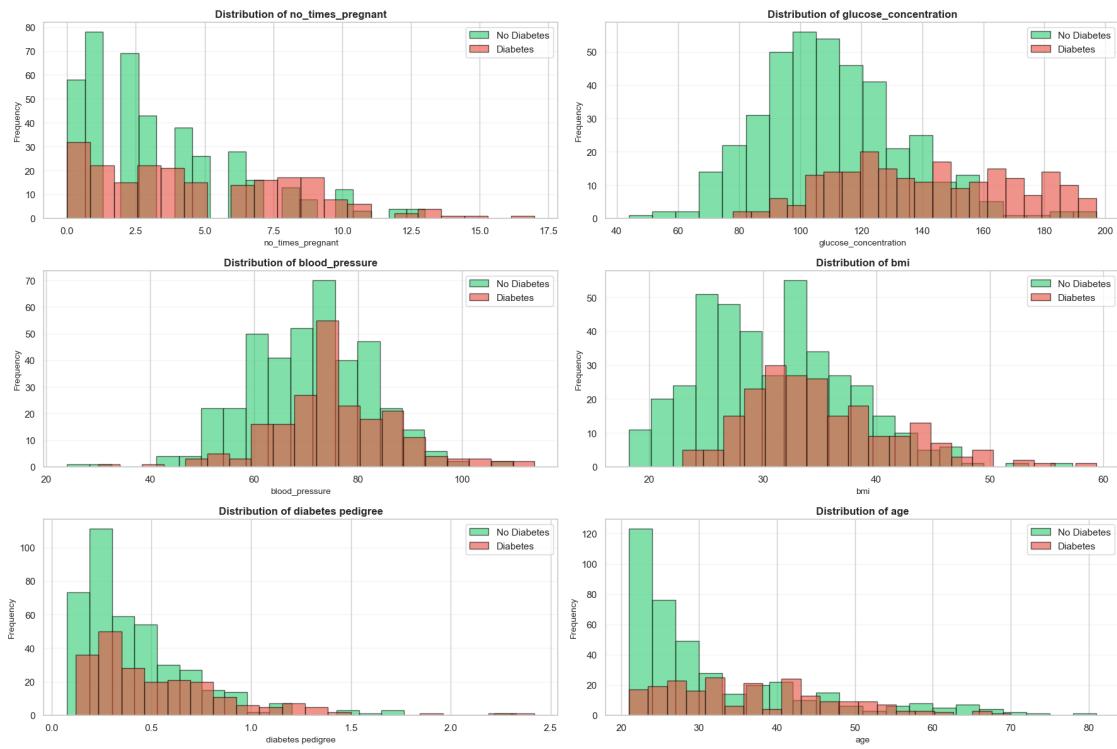
plt.tight_layout()
plt.show()

```

=====

FEATURE DISTRIBUTIONS BY DIABETES

=====



5.3.1 REMARKS

- Obvious difference in distributions of `glucose_concentration`, `bmi`, and `age` between diabetic and non-diabetic patients.
- Diabetic patients tend to have **higher glucose concentration, BMI, and age**.
- Other features (`no_times_pregnant`, `blood_pressure`, `diabetes pedigree`) have similar, overlapping distributions.

5.4 CORRELATION ANALYSIS

```
[18]: print_section('CORRELATION ANALYSIS')

### CALCULATE CORRELATION MATRIX
corr_matrix = train_clean[feature_cols_clean + ['diabetes']].corr()

### CREATE A FIGURE OF TWO SUBPLOTS (1 ROW X 2 COLS)
fig, axes = plt.subplots(1, 2)

### 1. CORRELATION HEATMAP IN LEFT SUBPLOT
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm',
             center=0, square=True, linewidths=1, cbar_kws={"shrink": 0.8}, ▾
             ax=axes[0])
axes[0].set_title('Feature Correlation Heatmap', fontsize=14, ▾
                  fontweight='bold', pad=20)

### 2. HORIZONTAL BAR PLOT IN RIGHT SUBPLOT
diabetes_corr = corr_matrix['diabetes'].round(3).drop('diabetes').sort_values()
# Dynamic color based on corr sign
colorsCorr = [colors[0] if x > 0 else colors[1] for x in diabetes_corr.values]
ax = diabetes_corr.plot(kind='barh', color=colorsCorr, ax=axes[1])

# Add corr values as text labels at the end of each bar
for i, (value, feature) in enumerate(zip(diabetes_corr.values, diabetes_corr. ▾
                                         index)):
    # x-position = end of bar (value)
    # y-position = bar center (i)
    ax.text(value + 0.01 if value > 0 else value - 0.01, # offset outside the ▾
            bar; dynamic based on corr sign
            i, # y-coordinate of horizontal bar
            f'{value:.3f}', # formatted corr value to 3 dp
            va='center',
            ha='left' if value > 0 else 'right') # outside bar; dynamic based ▾
            on corr sign

axes[1].set_title('Feature Correlations with Diabetes', fontsize=14, ▾
                  fontweight='bold')
axes[1].set_xlabel('Correlation Coefficient', fontsize=12)
axes[1].set_ylabel('Features', fontsize=12)

plt.tight_layout()
plt.show()
```

```
=====
CORRELATION ANALYSIS
=====
```



5.4.1 REMARKS

- Glucose concentration has highest linear correlation with diabetes (0.502).
- BMI is second most correlated with diabetes (0.297).
- Other features show moderate-to-low correlation.
- Some features are correlated with each other (multicollinearity), such as no_times_pregnant & age (0.53).

5.5 BOX PLOTS BY DIABETES

```
[19]: print_section('BOX PLOTS BY DIABETES STATUS')

### BUILD BOX PLOTS FOR ALL FEATURES BY DIABETES STATUS
n_cols = 3
n_rows = int(np.ceil(len(feature_cols_clean) / n_cols))
fig, axes = plt.subplots(n_rows, n_cols, figsize=(14, n_rows*4))
axes = axes.ravel()

for idx, col in enumerate(feature_cols_clean):
    train_clean.boxplot(column=col, by='diabetes', ax=axes[idx], ▾
    patch_artist=True, return_type='dict')
    axes[idx].set_title(f'{col} vs. diabetes status', fontsize=12, ▾
    fontweight='bold')
    axes[idx].set_xlabel('Diabetes (0=No, 1=Yes)', fontsize=10)
    axes[idx].set_ylabel(col, fontsize=10)
    axes[idx].grid(axis='y', alpha=0.3)

### HIDE UNUSED SUBPLOT
for idx in range(len(feature_cols_clean), len(axes)):
    axes[idx].axis('off')
```

```

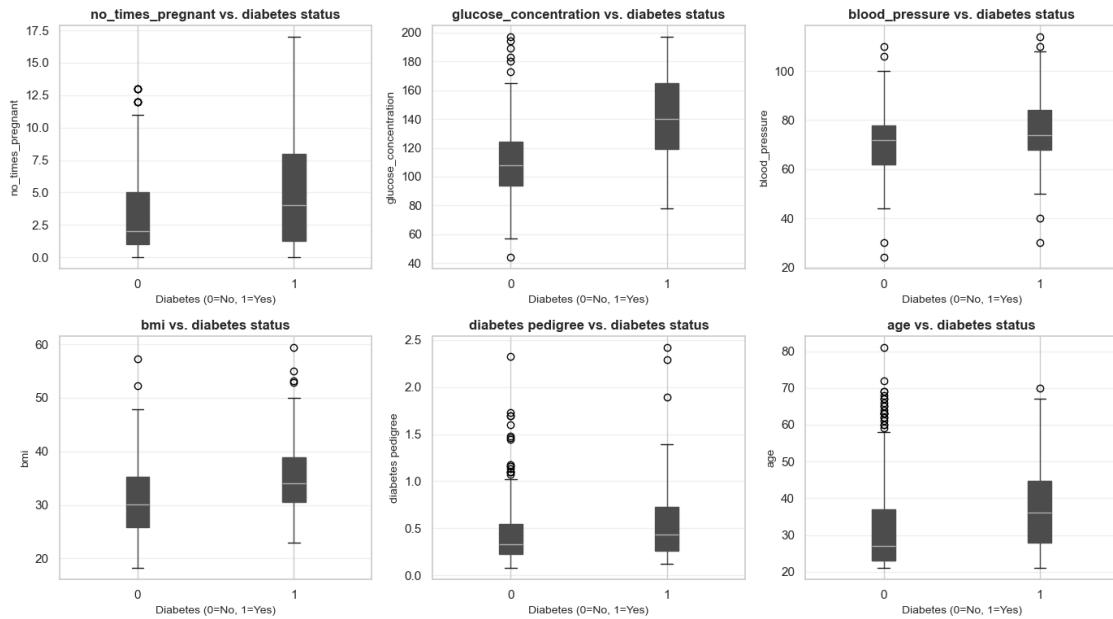
plt.suptitle('') # Remove default parent title
plt.tight_layout()
plt.show()

```

=====

BOX PLOTS BY DIABETES STATUS

=====



5.5.1 REMARKS

1. **Glucose concentration** shows the strongest separation between diabetic and non-diabetic groups.
 - Diabetic individuals have much higher glucose (higher median, higher quartiles, more extreme values).
 - This aligns with medical expectations and indicates glucose is the most discriminative feature.
2. Several features shift upward for diabetics (**age**, **BMI**, **pregnancies**).
 - **Age:** Diabetic individuals tend to be older on average.
 - **BMI:** Higher BMI among diabetics suggests a connection with obesity-related risk.
 - **Number of pregnancies:** The diabetic group shows higher median values & upper bound, but largely overlapping.
 - These features show moderate separation but not as strong as **glucose concentration**.
3. High overlap in **blood pressure** and **pedigree**.
 - Although the distributions overlap, diabetics tend to have:
 - Slightly higher blood pressure

- Higher diabetes pedigree function (genetic risk indicator)
- These features contribute information but *individually* are not determining factors to separate between statuses.

5.6 SCATTER PLOTS

```
[20]: print_section('SCATTER PLOTS')

### CHOOSE MOST IMPORTANT FEATURES ONLY
# glucose_concentration, bmi, age, diabetes pedigree

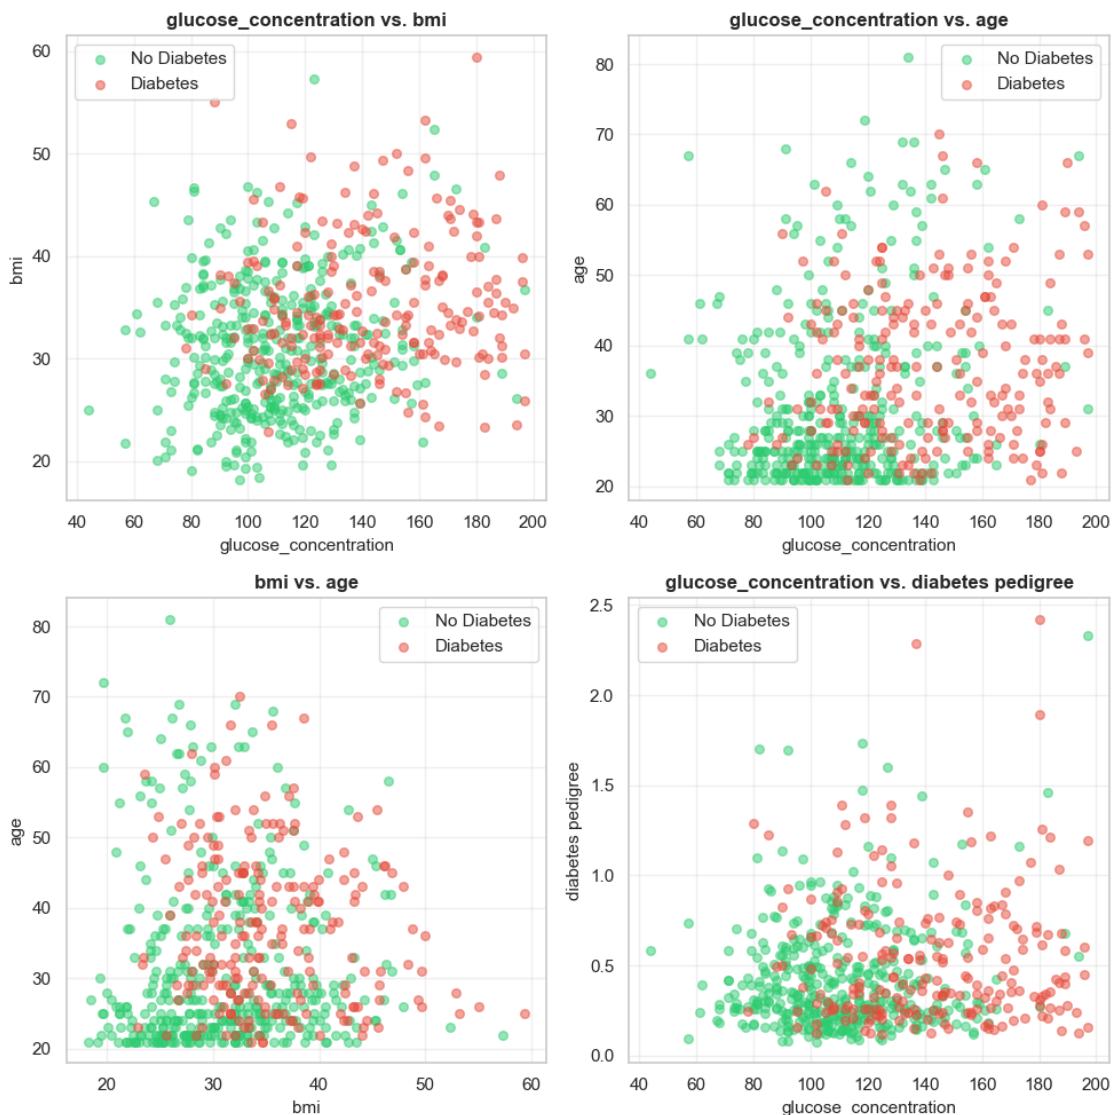
### CREATE A FIGURE OF 2x2 SUBPLOTS
fig, axes = plt.subplots(2, 2, figsize=(10, 10)) # Square plots for easier
    ↪visualization
axes = axes.ravel()

### DEFINE THE SCATTER PLOT COMBINATIONS
combinations = [
    ('glucose_concentration', 'bmi'),
    ('glucose_concentration', 'age'),
    ('bmi', 'age'),
    ('glucose_concentration', 'diabetes pedigree')
]
for idx, (feat1, feat2) in enumerate(combinations):
    ### PLOT NON-DIABETIC
    axes[idx].scatter(
        train_clean[train_clean['diabetes']==0][feat1],
        train_clean[train_clean['diabetes']==0][feat2],
        alpha=0.5, c='#2ecc71', label='No Diabetes', s=30
    )
    ### PLOT DIABETIC
    axes[idx].scatter(
        train_clean[train_clean['diabetes']==1][feat1],
        train_clean[train_clean['diabetes']==1][feat2],
        alpha=0.5, c='#e74c3c', label='Diabetes', s=30
    )
    ### SET PLOT ANNOTATIONS
    axes[idx].set_xlabel(feat1, fontsize=11)
    axes[idx].set_ylabel(feat2, fontsize=11)
    axes[idx].set_title(f'{feat1} vs. {feat2}', fontsize=12, fontweight='bold')
    axes[idx].legend()
    axes[idx].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```

=====

SCATTER PLOTS



5.6.1 REMARKS

- Glucose concentration is strongly associated with diabetes**
 - Across all plots involving glucose concentration (vs BMI, age, and diabetes pedigree), individuals with diabetes generally have higher glucose levels compared to those without diabetes.
 - This indicates that glucose concentration is a strong indicator of diabetes risk.
- BMI and age show moderate separation**
 - In the BMI vs age plot, individuals with diabetes tend to cluster at slightly higher BMI and older ages *compared to* those without diabetes, but there is considerable overlap.
 - This suggests BMI and age contribute to diabetes risk but are **not as definitive alone**.

3. Diabetes pedigree contributes to risk prediction

- In the glucose concentration vs diabetes pedigree plot, individuals with diabetes often have higher diabetes pedigree values, especially when combined with higher glucose levels.
- This shows that a family history of diabetes (captured by the diabetes pedigree) is an important factor when assessing diabetes risk.

Overall, glucose concentration appears to be the strongest single indicator, while BMI, age, and diabetes pedigree provide additional context that can improve risk assessment.

6 FEATURE ENGINEERING

```
[21]: train_before_FE = train_clean.copy()  
test_before_FE = test_clean.copy()
```

6.1 CREATE CATEGORICAL FEATURES

6.1.1 BMI CATEGORIES

- WHO classifications: <https://www.who.int/data/gho/data/themes/topics/topic-details/GHO/body-mass-index>

```
[22]: print_section('CREATE BMI CATEGORIES BASED ON WHO CLASSIFICATIONS')  
### DEFINE CATEGORICAL BINS AND LABELS BASED ON WHO CLASSIFICATIONS  
bins_bmi = [0, 18.5, 25, 30, float('inf')] # infinite upper bound  
labels_bmi = [0, 1, 2, 3] # 0=Underweight, 1=Normal, 2=Overweight, 3=Obese  
  
### CUT WITH LEFT-INCLUSIVE (DEFAULT MODE IS RIGHT INCLUSIVE right=True)  
train_clean['bmi_category'] = pd.cut(train_clean['bmi'], bins=bins_bmi,□  
    ↴labels=labels_bmi, right=False)  
test_clean['bmi_category'] = pd.cut(test_clean['bmi'], bins=bins_bmi,□  
    ↴labels=labels_bmi, right=False)  
  
print('BMI CATEGORIES CREATED:')  
print('0 = Underweight (<18.5)')  
print('1 = Normal (18.5-24.9)')  
print('2 = Overweight (25-29.9)')  
print('3 = Obese ( 30)')  
  
### SHOW CATEGORIES FOR TRAIN DATA  
print('\nDISTRIBUTION IN TRAIN DATA:')  
print(train_clean['bmi_category'].value_counts().sort_index())
```

=====

CREATE BMI CATEGORIES BASED ON WHO CLASSIFICATIONS

=====

BMI CATEGORIES CREATED:

0 = Underweight (<18.5)

1 = Normal (18.5-24.9)

```
2 = Overweight (25-29.9)
3 = Obese ( 30)
```

DISTRIBUTION IN TRAIN DATA:

```
bmi_category
0      2
1     78
2    155
3   379
Name: count, dtype: int64
```

6.1.2 AGE GROUPS

```
[23]: print_section('CREATE AGE GROUPS')
### DEFINE DESIRED CATEGORICAL BINS & LABELS
bins_age = [20, 30, 40, 50, float('inf')] # infinite upper bound
labels_age = [0, 1, 2, 3] # 0=20-29, 1=30-39, 2=40-49, 3=50+
### CUT WITH LEFT-INCLUSIVE (DEFAULT MODE IS RIGHT INCLUSIVE right=True)
train_clean['age_group'] = pd.cut(train_clean['age'], bins=bins_age, labels=labels_age, right=False)
test_clean['age_group'] = pd.cut(test_clean['age'], bins=bins_age, labels=labels_age, right=False)

print('AGE GROUPS CREATED:')
print('0 = 20-29 years')
print('1 = 30-39 years')
print('2 = 40-49 years')
print('3 = 50+ years')

### SHOW CATEGORIES FOR TRAIN DATA
print('\nDISTRIBUTION IN TRAIN DATA:')
print(train_clean['age_group'].value_counts().sort_index())
```

=====

CREATE AGE GROUPS

=====

AGE GROUPS CREATED:

```
0 = 20-29 years
1 = 30-39 years
2 = 40-49 years
3 = 50+ years
```

DISTRIBUTION IN TRAIN DATA:

```
age_group
0      319
1     128
2      95
```

```
3      72
Name: count, dtype: int64
```

6.1.3 GLUCOSE RISK CATEGORIES

- Bins are based on https://diabetesjournals.org/care/article/29/suppl_1/s43/23313/Diagnosis-and-Classification-of-Diabetes-Mellitus

```
[24]: print_section('CREATE GLUCOSE RISK CATEGORIES')
    ### DEFINE CATEGORICAL BINS & LABELS
    bins_glucose = [0, 100, 126, float('inf')]
    labels_glucose = [0, 1, 2] # 0=Normal, 1=Prediabetic, 2=Diabetic

    ### CUT WITH LEFT-INCLUSIVE (DEFAULT MODE IS RIGHT INCLUSIVE right=True)
    train_clean['glucose_category'] = pd.cut(train_clean['glucose_concentration'], ↴
        bins=bins_glucose, labels=labels_glucose, right=False)
    test_clean['glucose_category'] = pd.cut(test_clean['glucose_concentration'], ↴
        bins=bins_glucose, labels=labels_glucose, right=False)

    print('GLUCOSE RISK CATEGORIES CREATED:')
    print('0 = Normal (<100 mg/dL)')
    print('1 = Prediabetic (100-125.9 mg/dL)')
    print('2 = Diabetic range ( 126 mg/dL)')

    ### SHOW CATEGORIES FOR TRAIN DATA
    print('\nDISTRIBUTION IN TRAIN DATA:')
    print(train_clean['glucose_category'].value_counts().sort_index())
```

```
=====
CREATE GLUCOSE RISK CATEGORIES
=====

GLUCOSE RISK CATEGORIES CREATED:
0 = Normal (<100 mg/dL)
1 = Prediabetic (100-125.9 mg/dL)
2 = Diabetic range ( 126 mg/dL)

DISTRIBUTION IN TRAIN DATA:
glucose_category
0      151
1      229
2      234
Name: count, dtype: int64
```

6.1.4 CONVERT CATEGORICAL FEATURES TO NUMERIC

```
[25]: train_clean['bmi_category'] = train_clean['bmi_category'].astype(int)
test_clean['bmi_category'] = test_clean['bmi_category'].astype(int)
```

```

train_clean['age_group'] = train_clean['age_group'].astype(int)
test_clean['age_group'] = test_clean['age_group'].astype(int)

train_clean['glucose_category'] = train_clean['glucose_category'].astype(int)
test_clean['glucose_category'] = test_clean['glucose_category'].astype(int)

print(' CATEGORICAL FEATURES (CONVERTED TO NUMERIC) CREATED!')

```

CATEGORICAL FEATURES (CONVERTED TO NUMERIC) CREATED!

6.2 CREATE RATIO & INTERACTION FEATURES

6.2.1 GLUCOSE-BMI RATIO

```

[26]: ### CREATE
train_clean['glucose_bmi_ratio'] = train_clean['glucose_concentration']/_
    ↪(train_clean['bmi'] + 1e-6)
test_clean['glucose_bmi_ratio'] = test_clean['glucose_concentration']/_
    ↪(test_clean['bmi'] + 1e-6)

### SHOW MEAN & RANGE FOR TRAIN DATA
print_section('CREATE GLUCOSE-BMI RATIO')
print(f"Mean: {train_clean['glucose_bmi_ratio'].mean():.2f}")
print(f"Range: {train_clean['glucose_bmi_ratio'].min():.2f} to_
    ↪{train_clean['glucose_bmi_ratio'].max():.2f}")

```

=====

CREATE GLUCOSE-BMI RATIO

=====

Mean: 3.86

Range: 1.48 to 8.26

6.2.2 BMI-AGE INTERACTION

```

[27]: ### CREATE
train_clean['bmi_age_interaction'] = train_clean['bmi'] * train_clean['age']
test_clean['bmi_age_interaction'] = test_clean['bmi'] * test_clean['age']

### SHOW MEAN FOR TRAIN DATA
print_section('CREATE BMI-AGE INTERACTION')
print(f"Mean: {train_clean['bmi_age_interaction'].mean():.2f}")
print(f"Range: {train_clean['bmi_age_interaction'].min():.2f} to_
    ↪{train_clean['bmi_age_interaction'].max():.2f}")

```

=====

CREATE BMI-AGE INTERACTION

=====

Mean: 1082.14

Range: 382.20 to 2697.00

6.2.3 PREGNANCY-AGE INTERACTION

```
[28]: ### CREATE
train_clean['pregnancy_age_interaction'] = train_clean['no_times_pregnant'] *_
    ↪train_clean['age']
test_clean['pregnancy_age_interaction'] = test_clean['no_times_pregnant'] *_
    ↪test_clean['age']

### SHOW MEAN FOR TRAIN DATA
print_section('CREATE PREGNANCY-AGE INTERACTION')
print(f"Mean: {train_clean['pregnancy_age_interaction'].mean():.2f}")
print(f"Range: {train_clean['pregnancy_age_interaction'].min():.2f} to"
    ↪{train_clean['pregnancy_age_interaction'].max():.2f}")

print('\n Ratio and interaction features created successfully!')
```

=====

CREATE PREGNANCY-AGE INTERACTION

=====

Mean: 149.79
Range: 0.00 to 799.00

Ratio and interaction features created successfully!

6.3 CREATE HIGH-RISK FLAGS

6.3.1 HIGH GLUCOSE FLAG

```
[29]: # DEFINE THRESHOLD
glucose_threshold = 126 # mg/dL (diabetic range)

### CREATE FLAGS
train_clean['high_glucose_flag'] = (train_clean['glucose_concentration'] >=
    ↪glucose_threshold).astype(int)
test_clean['high_glucose_flag'] = (test_clean['glucose_concentration'] >=
    ↪glucose_threshold).astype(int)

### SHOW FLAG LOGIC & COUNT FOR TRAIN DATA
print_section('CREATE HGH GLUCOSE FLAG')
print(f'high_glucose_flag = 1 if glucose_concentration >= {glucose_threshold}u'
    ↪'mg/dL; 0 otherwise')
print(f'{train_clean["high_glucose_flag"].sum()} patients'
    ↪ f'{train_clean["high_glucose_flag"].mean()*100:.1f}%' flagged in TRAIN data')
```

=====

CREATE HGH GLUCOSE FLAG

=====

high_glucose_flag = 1 if glucose_concentration >= 126 mg/dL; 0 otherwise
234 patients (38.1%) flagged in TRAIN data

6.3.2 HIGH BMI FLAG

```
[30]: ### DEFINE THRESHOLD
bmi_threshold = 30 # Obese category

### CREATE FLAGS
train_clean['high_bmi_flag'] = (train_clean['bmi'] >= bmi_threshold).astype(int)
test_clean['high_bmi_flag'] = (test_clean['bmi'] >= bmi_threshold).astype(int)

### SHOW FLAG LOGIC & COUNT FOR TRAIN DATA
print_section('CREATE HIGH BMI FLAG')
print(f'high_bmi_flag = 1 if BMI >= {bmi_threshold}; 0 otherwise')
print(f'{train_clean["high_bmi_flag"].sum()} patients')
    ↪({train_clean["high_bmi_flag"].mean()*100:.1f}%) flagged in TRAIN data')
```

```
=====
CREATE HIGH BMI FLAG
=====
high_bmi_flag = 1 if BMI >= 30; 0 otherwise
379 patients (61.7%) flagged in TRAIN data
```

6.3.3 HIGH BLOOD PRESSURE FLAG

- Based on <https://www.mayoclinic.org/diseases-conditions/high-blood-pressure/in-depth/blood-pressure/art-20050982>

```
[31]: ### DEFINE THRESHOLD
bp_threshold = 80 # mmHg (hypertension stage 1)

### CREATE FLAGS
train_clean['high_bp_flag'] = (train_clean['blood_pressure'] >= bp_threshold).
    ↪astype(int)
test_clean['high_bp_flag'] = (test_clean['blood_pressure'] >= bp_threshold).
    ↪astype(int)

### SHOW FLAG LOGIC & COUNT FOR TRAIN DATA
print_section('CREATE HIGH BLOOD PRESSURE FLAG')
print(f'high_bp_flag = 1 if BP >= {bp_threshold} mmHg; 0 otherwise')
print(f'{train_clean["high_bp_flag"].sum()} patients')
    ↪({train_clean["high_bp_flag"].mean()*100:.1f}%) flagged in TRAIN data')
```

```
=====
CREATE HIGH BLOOD PRESSURE FLAG
=====
high_bp_flag = 1 if BP >= 80 mmHg; 0 otherwise
168 patients (27.4%) flagged in TRAIN data
```

6.3.4 HIGH AGE FLAG

```
[32]: ### DEFINE THRESHOLD
age_threshold = 40

### CREATE FLAGS
train_clean['high_age_flag'] = (train_clean['age'] >= age_threshold).astype(int)
test_clean['high_age_flag'] = (test_clean['age'] >= age_threshold).astype(int)

### SHOW FLAG LOGIC & COUNT FOR TRAIN DATA
print_section('CREATE HIGH AGE FLAG')
print(f'high_age_flag = 1 if age >= {age_threshold}; 0 otherwise')
print(f'{train_clean["high_age_flag"].sum()} patients')
    ↪({train_clean["high_age_flag"].mean()*100:.1f}%) flagged in TRAIN data')
```

```
=====
CREATE HIGH AGE FLAG
=====
high_age_flag = 1 if age >= 40; 0 otherwise
167 patients (27.2%) flagged in TRAIN data
```

6.3.5 COMBINED RISK SCORE

```
[33]: ### SUM THE HIGH-RISK FLAGS INTO RISK SCORES COLUMNS
train_clean['risk_score'] = (train_clean['high_glucose_flag'] +
                             train_clean['high_bmi_flag'] +
                             train_clean['high_bp_flag'] +
                             train_clean['high_age_flag'])

test_clean['risk_score'] = (test_clean['high_glucose_flag'] +
                           test_clean['high_bmi_flag'] +
                           test_clean['high_bp_flag'] +
                           test_clean['high_age_flag'])

### SHOW COMBINED RISK SCORE FOR TRAIN DATA
print_section('CREATE COMBINED RISK SCORE')
print('risk_score: Sum of all high-risk flags (0-4)')
print(f'Mean risk score for train: {train_clean["risk_score"].mean():.2f}')

print(f'\nDISTRIBUTION IN TRAIN DATA:')
print(train_clean['risk_score'].value_counts().sort_index())

print('\n High-risk flags created successfully!')
```

```
=====
CREATE COMBINED RISK SCORE
=====
risk_score: Sum of all high-risk flags (0-4)
Mean risk score for train: 1.54
```

```
DISTRIBUTION IN TRAIN DATA:
```

```
risk_score  
0      127  
1      187  
2      172  
3       95  
4       33  
Name: count, dtype: int64
```

```
High-risk flags created successfully!
```

6.4 CREATE POLYNOMIAL FEATURES

```
[34]: print_section('CREATE SQUARED FEATURES FOR NON-LINEAR RELATIONSHIPS')  
  
### GLUCOSE SQUARED  
train_clean['glucose_squared'] = train_clean['glucose_concentration'] ** 2  
test_clean['glucose_squared'] = test_clean['glucose_concentration'] ** 2  
print('glucose_squared: Glucose concentration squared')  
  
### BMI SQUARED  
train_clean['bmi_squared'] = train_clean['bmi'] ** 2  
test_clean['bmi_squared'] = test_clean['bmi'] ** 2  
print('bmi_squared: BMI squared')  
  
### AGE SQUARED  
train_clean['age_squared'] = train_clean['age'] ** 2  
test_clean['age_squared'] = test_clean['age'] ** 2  
print('age_squared: Age squared')  
  
print('\n Polynomial features created successfully!')
```

```
=====
```

```
CREATE SQUARED FEATURES FOR NON-LINEAR RELATIONSHIPS
```

```
=====
```

```
glucose_squared: Glucose concentration squared  
bmi_squared: BMI squared  
age_squared: Age squared
```

```
Polynomial features created successfully!
```

6.5 CREATE LOG-TRANSFORMED FEATURES

```
[35]: print_section('CREATE LOG-TRANSFORMED FEATURES FOR SKEWED FEATURES')
```

```
### DIABETES PEDIGREE  
train_clean['pedigree_log'] = np.log1p(train_clean['diabetes pedigree'])
```

```

test_clean['pedigree_log'] = np.log1p(test_clean['diabetes pedigree'])
print('pedigree_log: Log-transformed diabetes pedigree function')

print('\n Log-transformed features created successfully!')

```

```
=====
CREATE LOG-TRANSFORMED FEATURES FOR SKEWED FEATURES
=====
pedigree_log: Log-transformed diabetes pedigree function

Log-transformed features created successfully!
```

6.6 FEATURE ENGINEERING SUMMARY

```
[36]: print_section('FEATURE ENGINEERING SUMMARY')

### FEATURE BEFORE AND AFTER FEATURE ENGINEERING
original_features = list(train_before_FE.columns)
new_features = [col for col in train_clean.columns if col not in
                ↪original_features]

print(f'\n--- FEATURE COUNTS ---')
print(f'Original feature count: {len(original_features)}')
print(f'New features created: {len(new_features)}')
print(f'Total feature count: {len(train_clean.columns)}')

### NEW FEATURES
print(f'\n--- NEW FEATURES CREATED ---')
for i, feat in enumerate(new_features, 1):
    print(f'{i:2d}. {feat}')

### DISPLAY SAMPLE VALUES OF ENGINEERED FEATURES
print(f'\n--- SAMPLE VALUES OF ENGINEERED FEATURES ---')
display(train_clean[new_features].head())

print('\n FEATURE ENGINEERING COMPLETED!')
print('Ready to proceed with model training.\n')
```

```
=====
FEATURE ENGINEERING SUMMARY
=====
```

```

--- FEATURE COUNTS ---
Original feature count: 8
New features created: 15
Total feature count: 23

```

```
--- NEW FEATURES CREATED ---
```

```

1. bmi_category
2. age_group
3. glucose_category
4. glucose_bmi_ratio
5. bmi_age_interaction
6. pregnancy_age_interaction
7. high_glucose_flag
8. high_bmi_flag
9. high_bp_flag
10. high_age_flag
11. risk_score
12. glucose_squared
13. bmi_squared
14. age_squared
15. pedigree_log

```

--- SAMPLE VALUES OF ENGINEERED FEATURES ---

	bmi_category	age_group	glucose_category	glucose_bmi_ratio	\
0	3	0	1	3.284457	
1	3	3	2	3.907104	
2	3	0	0	2.447368	
3	3	2	1	3.084833	
4	3	0	2	3.950617	
	bmi_age_interaction	pregnancy_age_interaction	high_glucose_flag	\	
0	886.6		52	0	
1	1866.6		561	1	
2	874.0		46	0	
3	1594.9		41	0	
4	874.8		81	1	
	high_bmi_flag	high_bp_flag	high_age_flag	risk_score	glucose_squared \
0	1	0	0	1	12544
1	1	1	1	4	20449
2	1	0	0	1	8649
3	1	1	1	3	14400
4	1	0	0	2	16384
	bmi_squared	age_squared	pedigree_log		
0	1162.81	676	0.273837		
1	1339.56	2601	0.226338		
2	1444.00	529	0.515216		
3	1513.21	1681	0.771034		
4	1049.76	729	0.437610		

FEATURE ENGINEERING COMPLETED!

Ready to proceed with model training.

```
[37]: ### DISPLAY FINAL DATASETS READY FOR ML
print_section('FINAL DATASETS READY FOR ML')
print(f'Train shape: {train_clean.shape}')
print(f'Test shape: {test_clean.shape}')
print(f'\nFirst 5 rows of TRAIN data:')
display(train_clean.head())
```

=====

FINAL DATASETS READY FOR ML

=====

Train shape: (614, 23)

Test shape: (154, 22)

First 5 rows of TRAIN data:

	p_id	no_times_pregnant	glucose_concentration	blood_pressure	bmi	\
0	316	2		112	68	34.1
1	25	11		143	94	36.6
2	710	2		93	64	38.0
3	658	1		120	80	38.9
4	542	3		128	72	32.4

	diabetes	pedigree	age	diabetes	bmi_category	age_group	\
0		0.315	26	0	3	0	
1		0.254	51	1	3	3	
2		0.674	23	1	3	0	
3		1.162	41	0	3	2	
4		0.549	27	1	3	0	

	glucose_category	glucose_bmi_ratio	bmi_age_interaction	\
0	1	3.284457	886.6	
1	2	3.907104	1866.6	
2	0	2.447368	874.0	
3	1	3.084833	1594.9	
4	2	3.950617	874.8	

	pregnancy_age_interaction	high_glucose_flag	high_bmi_flag	high_bp_flag	\
0	52	0	1	0	
1	561	1	1	1	
2	46	0	1	0	
3	41	0	1	1	
4	81	1	1	0	

	high_age_flag	risk_score	glucose_squared	bmi_squared	age_squared	\
0	0	1	12544	1162.81	676	
1	1	4	20449	1339.56	2601	
2	0	1	8649	1444.00	529	

3	1	3	14400	1513.21	1681
4	0	2	16384	1049.76	729

```
pedigree_log
0    0.273837
1    0.226338
2    0.515216
3    0.771034
4    0.437610
```

7 MACHINE LEARNING

- Goal: Train classification models and select the best one.
- Algorithms: k-Nearest Neighbors, Decision Tree, Random Forest.

```
[38]: ### MAKE COPIES TO PRESERVE DATA
train_before_ML = train_clean.copy()
test_before_ML = test_clean.copy()
```

7.1 PREP DATA FOR MODELING

```
[39]: print_section('PREP DATA FOR MODELING')

### DROP NON-PREDICTIVE IDENTIFIER (p_id) AND SEPARATE FEATURES (X) FROM TARGET ↴(y)
X = train_clean.drop(['p_id', 'diabetes'], axis=1)
y = train_clean['diabetes']
X_test_final = test_clean.drop(['p_id'], axis=1) # Reserve for final predictions

### SPLIT TRAIN DATA INTO 80% TRAINING & 20% VALIDATION
# Set the random_state seed (42) to ensure same random split for every run
# stratify=y to have equal proportion of each class in train & val during ↴split, based on y
X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f'Train set: {X_train.shape}')
print(f'Validation set: {X_val.shape}')

### FEATURE SCALING
scaler = StandardScaler()
# FIT THE SCALER ON TRAIN SETS ONLY AND TRANSFORM IT
X_train_scaled = scaler.fit_transform(X_train)
# TRANSFORM THE VALIDATION AND TEST SETS USING THE SAME SCALER
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test_final)
```

```
print('\n Data prepared for modeling!')
```

```
=====
```

```
PREP DATA FOR MODELING
```

```
=====
```

```
Train set: (491, 21)  
Validation set: (123, 21)
```

```
Data prepared for modeling!
```

7.2 TRAIN CLASSIFICATION MODELS

```
[40]: ### INITIALIZE MODELS IN A DICTIONARY  
models = {  
    'k-Nearest Neighbors': KNeighborsClassifier(n_neighbors=5),  
    'Decision Tree': DecisionTreeClassifier(random_state=42),  
    'Random Forest': RandomForestClassifier(random_state=42, n_estimators=100)  
}  
  
### DEFINE AN EMPTY LIST TO STORE THE TRAINING RESULTS  
results = []  
  
### LOOP THROUGH EACH MODEL  
for name, model in models.items():  
    ### SHOW WHAT MODEL IS BEING TRAINED  
    print(f'\nTraining {name}...')  
  
### TRAIN MODEL  
model.fit(X_train_scaled, y_train)  
  
### USE TRAINED MODEL TO PREDICT THE TARGET (y) IN THE VALIDATION SET  
y_pred = model.predict(X_val_scaled)  
  
### CALCULATE METRICS  
accuracy = accuracy_score(y_val, y_pred) # % of correct predictions overall  
precision = precision_score(y_val, y_pred) # % of predicted positives that  
    ↪are actually positive  
recall = recall_score(y_val, y_pred) # % of actual positives correctly  
    ↪predicted  
f1 = f1_score(y_val, y_pred) # Harmonic mean of precision & recall  
    ↪(balances both)  
  
### STORE RESULTS  
results.append({  
    'Model': name,  
    'Accuracy': accuracy,
```

```

        'Precision': precision,
        'Recall': recall,
        'F1-Score': f1
    })

### SHOW REPORT
print(f' {name} trained')
print(f'- Accuracy: {accuracy:.3f}')
print(f'- Precision: {precision:.3f}')
print(f'- Recall: {recall:.3f}')
print(f'- F1-Score: {f1:.3f}')

print('\n All models trained successfully!')

```

Training k-Nearest Neighbors...

 k-Nearest Neighbors trained

- Accuracy: 0.756
- Precision: 0.659
- Recall: 0.628
- F1-Score: 0.643

Training Decision Tree...

 Decision Tree trained

- Accuracy: 0.748
- Precision: 0.625
- Recall: 0.698
- F1-Score: 0.659

Training Random Forest...

 Random Forest trained

- Accuracy: 0.821
- Precision: 0.733
- Recall: 0.767
- F1-Score: 0.750

All models trained successfully!

7.3 COMPARE MODEL PERFORMANCE

```
[41]: ### CREATE COMPARISON DF FROM RESULTS, SORTED BY DESCENDING F1-Score
results_df = pd.DataFrame(results)
results_df = results_df.sort_values('F1-Score', ascending=False)

print('--- MODEL PERFORMANCE METRICS COMPARISON ---')
display(results_df)

### VISUALIZE COMPARISON IN BAR CHART
```

```

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

### LEFT SUBPLOT axes[0] TO COMPARE ALL MODEL PERFORMANCE METRICS
# Plot on the left
ax1 = axes[0]
results_df.plot(x='Model', y=['Accuracy', 'Precision', 'Recall', 'F1-Score'],
                 kind='bar', ax=ax1)

# Set plot annotations
ax1.set_title('Model Performance Comparison', fontsize=14, fontweight='bold')
ax1.set_ylabel('Score', fontsize=12)
ax1.set_xlabel('')
ax1.set_xticklabels(ax1.get_xticklabels(), rotation=45, ha='right')
ax1.legend(loc='lower right')
ax1.grid(axis='y', alpha=0.3)
ax1.set_ylim([0, 1])

### RIGHT SUBPLOT axes[1] TO COMPARE F1-Score ONLY
# Plot on the right
ax2 = axes[1]
results_df_sorted = results_df.sort_values('F1-Score')

# Set plot annotations
ax2.banh(results_df_sorted['Model'], results_df_sorted['F1-Score'], color="#3498db")
ax2.set_title('F1-Score Ranking', fontsize=14, fontweight='bold')
ax2.set_xlabel('F1-Score', fontsize=12)
ax2.grid(axis='x', alpha=0.3)

# Add F1-Score value labels
for i, v in enumerate(results_df_sorted['F1-Score']):
    ax2.text(v + 0.01, i, f'{v:.3f}', va='center')

plt.tight_layout()
plt.show()

# Identify best model
best_model_name = results_df.iloc[0]['Model']
best_f1 = results_df.iloc[0]['F1-Score']

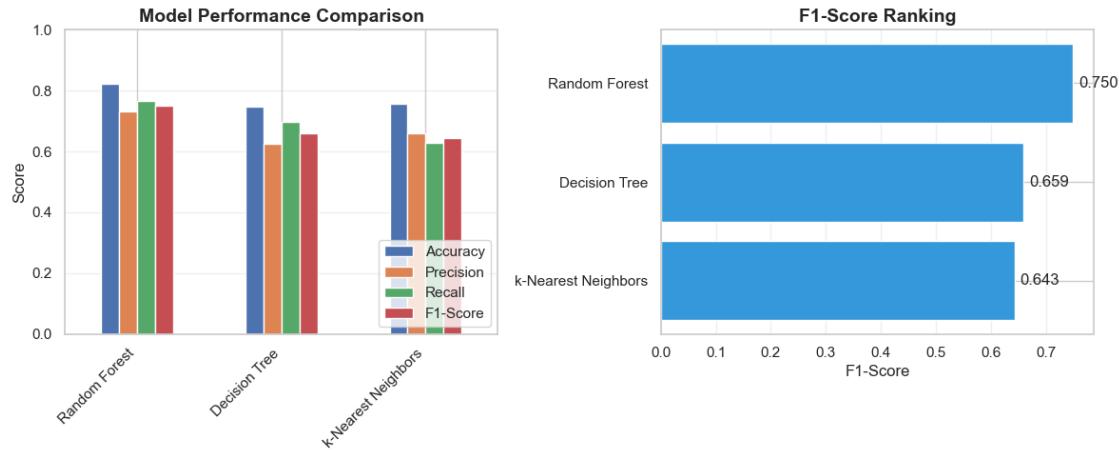
print(f'\n BEST MODEL: {best_model_name}')
print(f'    F1-Score: {best_f1:.3f}')

```

--- MODEL PERFORMANCE METRICS COMPARISON ---

	Model	Accuracy	Precision	Recall	F1-Score
2	Random Forest	0.821138	0.733333	0.767442	0.750000
1	Decision Tree	0.747967	0.625000	0.697674	0.659341

```
0 k-Nearest Neighbors 0.756098 0.658537 0.627907 0.642857
```



BEST MODEL: Random Forest

F1-Score: 0.750

7.4 EVALUATE BEST MODEL

```
[42]: ### GET BEST MODEL AND ITS PREDICTIONS ON y IN THE VALIDATION SET
best_model = models[best_model_name]
y_pred_best = best_model.predict(X_val_scaled)

### CONFUSION MATRIX TO SEE WHERE THE MODEL IS MAKING MISTAKES
print(f'--- CONFUSION MATRIX: {best_model_name} ---')
cm = confusion_matrix(y_val, y_pred_best)

### VISUALIZE CONFUSION MATRIX
plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', # annot=True adds the values;
            fmt='d' present them as int
            xticklabels=['No Diabetes', 'Diabetes'],
            yticklabels=['No Diabetes', 'Diabetes'])

# Set plot annotations
plt.title(f'Confusion Matrix - {best_model_name}', fontsize=14,
          fontweight='bold')
plt.ylabel('True Label', fontsize=12)
plt.xlabel('Predicted Label', fontsize=12)
plt.tight_layout()
plt.show()

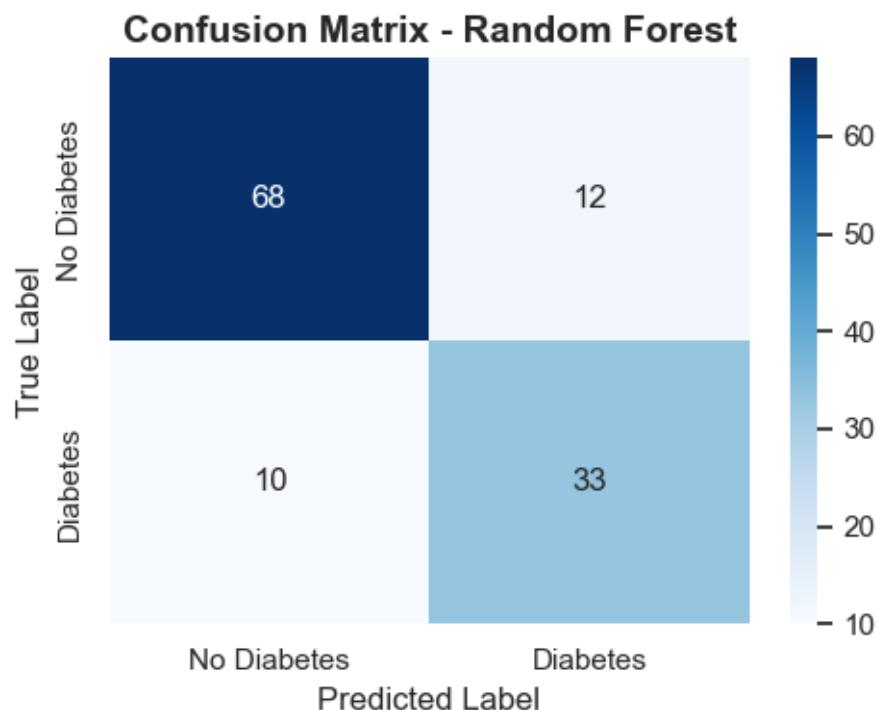
### CLASSIFICATION REPORT
```

```

# Precision: Of all patients predicted as diabetics, how many actually are?
# Recall: Of all actual diabetics, how many were correctly identified?
# F1-Score: Balances precision and recall - important when false negatives are
#           costly.
# Support: Number of true samples for each class, gives context to the metrics
print(f'\n--- Classification Report: {best_model_name} ---')
print(classification_report(y_val, y_pred_best, target_names=['No Diabetes',
#           'Diabetes']))

```

--- CONFUSION MATRIX: Random Forest ---



--- Classification Report: Random Forest ---

	precision	recall	f1-score	support
No Diabetes	0.87	0.85	0.86	80
Diabetes	0.73	0.77	0.75	43
accuracy			0.82	123
macro avg	0.80	0.81	0.81	123
weighted avg	0.82	0.82	0.82	123

CONFUSION MATRIX - Most non-diabetic patients are correctly identified (68/80). - Some diabetic patients are missed (10/43).

CLASSIFICATION REPORT - Accuracy: 0.81 => 82% of predictions are correct overall.
- Macro avg (unweighted): 0.81 => average of F1 across classes (diabetics & non-diabetics), treats both classes equally.
- Weighted avg: 0.82 => average of F1 weighted by class size, slightly favors the majority class (non-diabetics).
- F1-Score for diabetic patients (0.75) is lower than for non-diabetic (0.86) => the model struggles more with detecting diabetics.

7.5 MAKE PREDICTIONS ON TEST DATASET USING BEST MODEL

```
[43]: ### MAKE PREDICTIONS ON TEST DATASET USING BEST MODEL
test_predictions = best_model.predict(X_test_scaled)

### CREATE A DF OF PREDICTION RESULTS
predictions_df = pd.DataFrame({
    'p_id': test_clean['p_id'],
    'predicted_diabetes': test_predictions
})

print(f" Predictions completed for {len(predictions_df)} test samples")
print('Prediction distribution:')
print(predictions_df['predicted_diabetes'].value_counts())

### SHOW SAMPLE PREDICTIONS
print('\nSample predictions:')
display(predictions_df.head())
```

```
Predictions completed for 154 test samples
Prediction distribution:
predicted_diabetes
0      97
1      57
Name: count, dtype: int64
```

Sample predictions:

```
p_id  predicted_diabetes
0    437              1
1    411              0
2    639              0
3    213              1
4    181              0
```

7.6 ML SUMMARY

```
[44]: print_section('ML RESULTS SUMMARY')
print(f'''- ALGORITHMS TESTED: {', '.join(models)}
- BEST MODEL: {best_model_name}
- PERFORMANCE (VALIDATION SET):
  - Accuracy: {results_df.iloc[0]["Accuracy"]:.4f}
  - Precision: {results_df.iloc[0]["Precision"]:.4f}'''
```

```

    - Recall:    {results_df.iloc[0]["Recall"]:.4f}
    - F1-Score:  {results_df.iloc[0]["F1-Score"]:.4f}
- TEST PREDICTIONS:
    - Total samples: {len(predictions_df)}
    - Predictions saved in: predictions_df''')

print('\n ML COMPLETED!')
print(' PROJECT COMPLETED!')

```

```
=====
ML RESULTS SUMMARY
=====
- ALGORITHMS TESTED: k-Nearest Neighbors, Decision Tree, Random Forest
- BEST MODEL: Random Forest
- PERFORMANCE (VALIDATION SET):
    - Accuracy:  0.8211
    - Precision: 0.7333
    - Recall:    0.7674
    - F1-Score:  0.7500
- TEST PREDICTIONS:
    - Total samples: 154
    - Predictions saved in: predictions_df
```

ML COMPLETED!

PROJECT COMPLETED!